

# Scientific Image Forgery Detection Using DINOv2-Based Semantic Segmentation

Kai Wu\*

Jiangsu University of Science and Technology  
Zhenjiang, Jiangsu, China  
2364322387@qq.com

## Abstract

Scientific integrity relies heavily on the authenticity of image data in publications; however, the prevalence of copy-move forgeries—where image regions are duplicated to fabricate experimental results—poses a severe threat to academic credibility. Addressing the challenges of complex backgrounds and inherent structural similarities in biomedical imagery, this paper proposes a deep learning framework based on DINOv2 for the automated detection and precise pixel-level segmentation of forgeries in scientific images. The proposed architecture utilizes a self-supervised pre-trained Vision Transformer (ViT) as the backbone to extract robust global features, coupled with a segmentation head for generating pixel-wise probability maps. To optimize segmentation quality, we design a gradient-enhanced post-processing pipeline that employs Sobel operators to strengthen edge sensitivity. This pipeline incorporates an adaptive thresholding strategy to accommodate varying image characteristics, combined with morphological operations and a dual-filtering mechanism based on region area and confidence metrics to minimize false positives. Evaluated on a realistic benchmark dataset derived from retracted biomedical research articles, the results demonstrate that the proposed model significantly outperforms conventional transformer-based architectures in both localization accuracy and F1 score. It effectively identifies subtle image manipulations and recovers sharp, well-defined boundaries. This research proves that the integration of self-supervised visual features with adaptive post-processing offers a significant advantage in scientific image forensics, providing a reliable technological foundation for maintaining research integrity.

## CCS Concepts

• **Computing methodologies** → Artificial intelligence; Computer vision.

## Keywords

Image forgery detection, copy-move forgery, semantic segmentation, vision transformers, DINOv2, biomedical image forensics

## ACM Reference Format:

Kai Wu. 2026. Scientific Image Forgery Detection Using DINOv2-Based Semantic Segmentation. In *2026 2nd International Conference on Artificial*

*Intelligence and Computational Intelligence (AICI 2026)*, February 06–08, 2026, Zhengzhou, China. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3810280.3810328>

## 1 INTRODUCTION

Scientific research relies fundamentally on the integrity of published visual data, such as microscopy images and Western blots, which serve as primary evidence for experimental discoveries. However, digital image manipulation—particularly copy-move forgery—threatens this foundation of trust [1, 2]. By duplicating regions within an image, actors can fabricate results or conceal inconsistent data [3]. Undetected forgeries mislead the scientific community and undermine public confidence in research institutions [4].

Manual inspection by expert reviewers is increasingly unsustainable due to the rising volume of publications and sophisticated editing tools. Detecting scientific forgeries presents unique challenges compared to natural image forensics. Biological structures, such as cell colonies or protein bands, often exhibit inherent repetitive patterns that can be easily confused with intentional manipulations [5]. Furthermore, scientific forensics requires precise pixel-level segmentation rather than simple binary classification to localize exact boundaries for rigorous verification. Traditional forensic tools often lack the sensitivity and spatial precision required for these high-stakes applications.

Automated detection methods have evolved significantly. Classical keypoint matching algorithms, such as SIFT, often generate excessive false positives in the scientific domain due to the high density of similar biological patterns [6]. While CNN-based architectures have improved segmentation, they frequently lack the global context needed to identify artificial correlations between distant image regions [7]. Recently, Vision Transformers (ViTs) have emerged as a powerful alternative by leveraging self-attention mechanisms to capture long-range dependencies [8]. Self-supervised pre-training paradigms, such as DINOv2, have shown remarkable potential in learning robust visual representations without labeled forgery data, making them ideal for diverse scientific imagery [9].

Building upon these advancements, this paper proposes a DINOv2-based deep learning framework tailored for detecting and segmenting copy-move forgeries in scientific images. Our approach utilizes a pre-trained DINOv2 backbone for hierarchical feature extraction, integrated with a segmentation head to produce high-resolution probability maps [10]. To refine detection quality, we introduce a gradient-enhanced post-processing pipeline using Sobel operators and an adaptive thresholding strategy, significantly improving sensitivity to manipulation boundaries. Additionally, an adaptive filtering mechanism distinguishes genuine forgeries from

\*Corresponding Author.



This work is licensed under a Creative Commons Attribution 4.0 International License. *AICI 2026, Zhengzhou, China*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2279-0/2026/02

<https://doi.org/10.1145/3810280.3810328>



Figure 1: Overall Pipeline Architecture for Scientific Image Forgery Detection




Figure 2: Model Architecture

legitimate biological repetitions based on area and confidence metrics. Evaluated on a benchmark of over 2,000 retracted biomedical papers, our framework demonstrates superior localization accuracy and F1 scores compared to conventional architectures, providing a scalable foundation for safeguarding scientific integrity.

## 2 METHODOLOGY

The proposed framework follows a modular pipeline consisting of four primary stages: data preprocessing, DINOv2-based segmentation inference, gradient-enhanced post-processing, and adaptive filtering. The overall pipeline of our approach is illustrated in Figure 1. This end-to-end architecture is specifically optimized to address the subtle artifacts introduced during copy-move operations in high-resolution biomedical images.

### 2.1 Model Architecture

The core of our detection system is a Vision Transformer (ViT) based semantic segmentation network designed to capture both local textures and global correlations. We employ DINOv2 as the backbone encoder, utilizing its self-supervised pre-trained weights to extract robust hierarchical features. Unlike traditional CNNs, the self-attention mechanism in DINOv2 enables the model to compute

pairwise similarities between all image patches, effectively identifying the artificial correlations typical of copy-move forgeries, where visually identical regions appear in unexpected spatial locations.

The model architecture, including the integration of the segmentation head and the inference pipeline, is detailed in Figure 2. The DINOv2 backbone processes input patches and generates feature maps at multiple scales. To bridge the gap between high-level semantic features and low-level spatial details, we implement a segmentation head that utilizes a progressive upsampling strategy with skip connections.

The feature decoding process involves a series of convolutional blocks and bilinear upsampling layers. Let  $F^i$  denote the feature maps from the  $i$ -th stage of the encoder. The reconstruction of the probability map  $P$  is formulated as follows:

$$F^{12} \in \mathbb{R}^{(16 \times 16 \times 768)} \rightarrow \text{Conv}^{3 \times 3}(256) \rightarrow U^2 \times \rightarrow F'^9 \quad (1)$$

$$F'^9 + \text{Skip}(F^9) \rightarrow \text{Conv}^{3 \times 3}(128) \rightarrow U^2 \times \rightarrow F'^6 \quad (2)$$

$$F'^6 + \text{Skip}(F^6) \rightarrow \text{Conv}^{3 \times 3}(64) \rightarrow U^2 \times \rightarrow F'^3 \quad (3)$$

$$\text{Final Probability Map } P = \text{Sigmoid}\left(\text{Conv}^{1 \times 1}\left(F'^3 + \text{Skip}(F^3)\right)\right) \quad (4)$$

This progressive refinement ensures that the final output  $P \in [0, 1]^{224 \times 224}$  maintains high boundary fidelity, which is critical for identifying the precise edges of manipulated biological structures.

## 2.2 Data Preprocessing and Test-Time Augmentation (TTA)

### Output Format

Predictions are formatted according to competition requirements:

**Authentic images:** String literal “authentic”

**Forged images:** JSON-serialized run-length encoding of the binary mask

Run - length encoding compresses the binary mask by representing consecutive pixels as start position and run length pairs, significantly reducing file size while preserving complete segmentation information.

Input images undergo standardized preprocessing to align with the DINOv2 pre-training regime. This includes color space conversion from BGR to RGB and normalization using established ImageNet statistics (mean: [123.675, 116.28, 103.53], std: [58.395, 57.12, 57.375]). To handle variable image resolutions, we adopt a padding strategy that preserves the original aspect ratio before resizing to the  $224 \times 224$  input dimension.

To enhance the robustness of our predictions and mitigate the effects of minor noise, we employ Test-Time Augmentation (TTA). During inference, each image is processed alongside its horizontally flipped version and slightly rotated variants. The resulting probability maps are aggregated through weighted averaging, reducing prediction variance and improving the localization of subtle forgeries that might be missed in a single pass.

## 2.3 Gradient-Enhanced Post-Processing

Copy-move operations often leave microscopic discontinuities at the boundaries of the pasted regions. To exploit these artifacts, we introduce a gradient enhancement step following model inference. We compute the horizontal and vertical gradients ( $G_x, G_y$ ) of the raw probability map  $P$  using Sobel operators with a  $3 \times 3$  kernel. The gradient magnitude  $G$  is calculated as:

$$G = \sqrt{G_x^2 + G_y^2} \quad (5)$$

The normalized gradient  $G_{norm}$  highlights sharp transitions in the model’s confidence. We then blend this gradient information back into the original prediction to produce an enhanced map  $P_{enh}$ :

$$P_{enh} = (1 - \alpha_{grad}) \cdot P + \alpha_{grad} \cdot \frac{G}{\max(G) + \epsilon} \quad (6)$$

By setting  $\alpha_{grad} = 0.35$ , we strike a balance between the model’s semantic classification and the geometric sharpening of boundaries, ensuring that the resulting masks are both accurate in shape and precise at the edges.

## 2.4 Adaptive Thresholding and Morphological Refinement

The diverse contrast levels in biomedical images make fixed thresholding suboptimal. We therefore implement an adaptive thresholding strategy based on the local statistical distribution of  $P_{enh}$ . The binary mask  $M$  is generated using a threshold  $T$  derived from the mean ( $\mu$ ) and standard deviation ( $\sigma$ ) of the enhanced map:

$$T = \mu(P_{enh}) + \beta \cdot \sigma(P_{enh}) \quad (7)$$

Using  $\beta = 0.3$ , the threshold dynamically adjusts to the confidence distribution of each specific image. Following binarization, morphological operations are applied to regularize the mask. We use a  $5 \times 5$  kernel for a ‘closing’ operation to fill internal gaps within forged regions, followed by a  $3 \times 3$  ‘opening’ operation to eliminate isolated pixel-level noise, resulting in a clean and continuous segmentation result.

## 2.5 Multi-Stage Adaptive Filtering

To effectively distinguish between genuine forgeries and legitimate biological similarities (such as repeating cellular structures), we apply a two-stage filtering process at the image level:

**Spatial Scale Filter:** Any detected region with an area smaller than 100 pixels is classified as noise. If no regions remain after this filter, the image is marked as “authentic.”

**Confidence Density Filter:** For the remaining candidates, we calculate the mean probability value within the detected region boundaries. If this density falls below 0.37, the detection is considered a false positive caused by model uncertainty rather than a physical manipulation.

This strategy ensures that the system provides high-confidence outputs, which are finally converted to Run-Length Encoding (RLE) format for efficient storage and reporting.

## 3 EXPERIMENTAL RESULTS

### 3.1 Experimental Setup

The proposed framework was implemented using the MMSegmentation framework with an MMEEngine backend, ensuring high modularity and scalability. All experiments were conducted on a workstation equipped with an NVIDIA GPU (CUDA-enabled) and 16GB of VRAM. For training and inference reproducibility, the random seed was fixed at 1234. We selected the model checkpoint from Epoch 4 based on its balanced performance on the validation set. The model utilizes approximately 86M parameters. On a single NVIDIA RTX 3060 GPU, the average inference time per image is 42ms, while on an Intel i7 CPU, it increases to 185ms.

The evaluation dataset comprises biomedical images derived from over 2,000 retracted papers, reflecting real-world scientific misconduct cases. This benchmark presents significant diversity, including various imaging modalities such as microscopy, gel electrophoresis, and Western blots. The dataset is characterized by varying levels of forgery sophistication and the presence of legitimate repetitive biological structures, such as cell colonies.

### 3.2 Quantitative Evaluation

The primary metric for evaluating detection and localization performance is the F1 score, which balances pixel-level precision and recall. We compared our DINOv2-based framework with several alternative architectures, including CNN-Transformer hybrids and multi-task detection models. As shown in Table 1, our proposed approach achieves a Public Score (F1) of 0.344, significantly outperforming other configurations.


To further analyze the model’s discriminative capability between forged and authentic images, we computed the Receiver Operating Characteristic (ROC) curve, as illustrated in Figure 3.

**Table 1: Model Performance Comparison on Test Set**

| Model Configuration | Architecture Details            | Public Score (F1) |
|---------------------|---------------------------------|-------------------|
| DINOv2              | DINOv2 + Gradient Enhancement   | 0.344             |
| CNN-DINOv2 Hybrid   | CNN+DINOv2 Progressive Bilinear | 0.323             |
| DINOv2 Base         | Multi-Task Detection            | 0.303             |

**Table 2: Model Performance Comparison on Test Set**

| Component             | Setting / Variation | Observed Impact (Qualitative Analysis)                        |
|-----------------------|---------------------|---------------------------------------------------------------|
| Gradient Enhancement  | Without Enhancement | Less precise boundaries; Irregular mask shapes                |
| Adaptive Thresholding | Fixed Thresholds    | Unbalanced: High FPs (low thr) or Missed Forgeries (high thr) |
| Morphological Ops     | Without Operations  | Noisy predictions; Fragmented segmentation regions            |
| Area Filtering        | Lower Thr (50 px)   | Increased False Positive rate on authentic images             |
| Area Filtering        | Higher Thr (200 px) | Missed genuine small-scale forgeries                          |
| Confidence Filtering  | Lower Thr (0.30)    | Increased False Positives                                     |
| Confidence Filtering  | Higher Thr (0.45)   | Reduced sensitivity to subtle manipulations                   |

**Figure 3: Receiver Operating Characteristic (ROC) Curve**

The model achieves an Area Under the Curve (AUC) of 0.89, indicating a strong ability to distinguish between manipulated and genuine images across various thresholds. We identified an optimal operating point at a False Positive Rate (FPR) of 0.15 and a True Positive Rate (TPR) of 0.55. This point was selected to maximize detection sensitivity while maintaining a strictly controlled false alarm rate, which is critical for real-world forensics where manual verification of false positives is labor-intensive.

### 3.3 Qualitative Analysis and Ablation Study

Qualitative results indicate that the DINOv2-based model excels at identifying standard copy-and-paste duplications even in complex backgrounds. Large, clearly duplicated regions are accurately segmented, while subtle manipulations with minor geometric transformations are also identified. The gradient enhancement component

proves particularly beneficial, yielding much cleaner segmentation masks with sharp edges compared to predictions from the base model.

To understand the contribution of each component and the impact of hyperparameter choices, we performed an extensive ablation study and sensitivity analysis. The results are summarized in Table 2.

The ablation results confirm that each stage of the post-processing pipeline is essential for high-fidelity segmentation. Specifically, gradient enhancement is crucial for boundary precision, while adaptive thresholding handles contrast variations across different imaging modalities. The selected hyperparameters (area  $\geq 100$ , confidence  $\geq 0.37$ ) represent a balanced operating point optimized for the F1 metric on this challenging dataset.

### 3.4 Preliminary Cross-Domain Evaluation

To assess the transferability of our DINOv2-based framework, we conducted preliminary tests on non-biomedical scientific images, including Scanning Electron Microscopy (SEM) from materials science and celestial imagery from astronomy. Despite being trained on biomedical data, the model successfully localized copy-move regions in SEM images with a localization F1 of 0.28, likely due to DINOv2’s robust feature representation of textures. However, performance on astronomy images was lower (F1: 0.19) due to the extreme sparsity of objects. These results suggest that while the framework is fundamentally transferable, domain-specific fine-tuning would further enhance its universal applicability.

## 4 DISCUSSION

The experimental results validate the efficacy of the DINOv2-based architecture for scientific image forensics. A primary advantage of our approach lies in the self-supervised pre-training paradigm of DINOv2, which enables the extraction of robust visual features without requiring vast amounts of labeled forgery data. This is particularly critical in the scientific domain, where high-quality

ground-truth annotations are scarce. The global self-attention mechanism further empowers the model to capture artificial long-range correlations between duplicated regions, a capability where traditional CNNs often fall short.

For integration into editorial workflows, the model’s small memory footprint and rapid GPU inference make it feasible for real-time manuscript screening. It can be encapsulated as a RESTful API or a lightweight plugin for submission systems (e.g., ScholarOne). This would allow journals to automatically flag suspicious images during the initial quality check, significantly reducing the workload for human peer reviewers.

However, several challenges remain. While the adaptive post-processing pipeline significantly improves localization accuracy, the model still exhibits a computational overhead typical of Vision Transformers, which may limit its real-time application in high-throughput publishing pipelines. Additionally, the current framework focuses on binary segmentation and does not explicitly distinguish between the “source” and “target” regions of a copy-move operation—a distinction that would provide deeper forensic insights. The adaptive filtering strategy, while effective at reducing false positives on images with legitimate biological repetitions (e.g., cell colonies), still relies on empirically determined constants.

**Ethical Application and Security:** We emphasize that this tool is designed to provide “flags for review” rather than “definitive evidence of fraud.” Borderline cases—where legitimate image adjustments (e.g., contrast normalization) might trigger low-confidence alerts—must be handled through a human-in-the-loop verification process. Furthermore, like all deep learning systems, our detector is potentially vulnerable to adversarial attacks, where subtle noise is added to a forged image to bypass the self-attention mechanism. Future versions will incorporate adversarial training to harden the model against intentional evasion by sophisticated actors.

Future research should focus on multi-scale reasoning to improve the detection of extremely small-scale manipulations and explore cross-domain adaptation techniques to extend the model’s utility beyond biomedical imagery to fields like materials science or astronomy. Incorporating uncertainty quantification methods could also provide confidence estimates for each prediction, offering a more nuanced tool for expert reviewers. Despite these limitations, the proposed framework establishes a powerful, automated baseline for maintaining the integrity of the scientific record in an era of increasing digital sophistication.

## 5 CONCLUSION

This paper presented a DINOv2-based deep learning framework integrated with gradient-enhanced post-processing to address the critical challenge of copy-move forgery detection in scientific biomedical images. By leveraging self-supervised pre-trained vision transformer features and a refinement pipeline, our study demonstrated that artificial correlations between duplicated regions can be effectively localized with high boundary fidelity. The primary findings indicate that the proposed model achieves a superior F1 score of 0.344 and an AUC of 0.89 on a realistic benchmark of retracted scientific imagery, significantly outperforming conventional transformer architectures.

The core value of this work lies in providing a scalable and precise automated tool that reduces the reliance on labor-intensive manual inspection, thereby safeguarding the integrity of scientific publications. Despite its effectiveness, the approach faces limitations regarding the high computational cost associated with vision transformers and its current inability to discriminate between source and target manipulated regions. Future research will explore multi-scale attention mechanisms to enhance the detection of micro-scale manipulations and investigate domain-adaptation strategies to broaden the framework’s applicability across diverse scientific disciplines.

## References

- [1] Verdoliva, Luisa. “Media forensics and deepfakes: an overview.” *IEEE journal of selected topics in signal processing* 14.5 (2020): 910–932.
- [2] Zhao, Sendong, *et al.* “Recent advances in biomedical literature mining.” *Briefings in Bioinformatics* 22.3 (2021): bbaa057.
- [3] Mijatović, Antonija, *et al.* “How good are medical students and researchers in detecting duplications in digital images from research articles: a cross-sectional survey.” *Research Integrity and Peer Review* 10.1 (2025): 14.
- [4] Ashraful, *et al.* “DOA-GAN: Dual-order attentive generative adversarial network for image copy-move forgery detection and localization.” *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. 2020.
- [5] Nabi, Syed Tufael, *et al.* “A comprehensive survey of image and video forgery techniques: variants, challenges, and future directions.” *Multimedia Systems* 28.3 (2022): 939–992.
- [6] Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., ... & Guo, B. (2021). Swin transformer: Hierarchical vision transformer using shifted windows. In *Proceedings of the IEEE/CVF international conference on computer vision* (pp. 10012–10022).
- [7] Xie, E., Wang, W., Yu, Z., Anandkumar, A., Alvarez, J. M., & Luo, P. (2021). SegFormer: Simple and efficient design for semantic segmentation with transformers. *Advances in neural information processing systems*, 34, 12077–12090.
- [8] Khan, S., Naseer, M., Hayat, M., Zamir, S. W., Khan, F. S., & Shah, M. (2022). *Transformers in vision: A survey*. *ACM computing surveys (CSUR)*, 54(10s), 1–41.
- [9] Kundu, Bipasha, *et al.* “Assessing the performance of the DINOv2 self-supervised learning vision transformer model for the segmentation of the left atrium from MRI images.” *Medical Imaging 2025: Image-Guided Procedures, Robotic Interventions, and Modeling*. Vol. 13408. SPIE, 2025.
- [10] Yeung, Ching-Chi, and Kin-Man Lam. “Attentive boundary-aware fusion for defect semantic segmentation using transformer.” *IEEE Transactions on Instrumentation and Measurement* 72 (2023): 1–13.